Finite task, action and horizon support

$$D_i^*(s) = J^*(s, g, U_0) - J^*(s, g, U_i), \qquad h_\ell(s, v) = E_{1:\ell}(\mathcal{R}(s, v))$$

A Exact sufficiency

A readout exists on the finite set iff identical h implies identical D*.

B Aliasing bound

$\max(e_1, e_2) \geq (\Delta - K\delta)_+ / 2$

K must be justified; full hidden state.

C Selection bound

$|\hat{J} - J^*| \leq \varepsilon$

Same units + same support:
regret $\leq 2\varepsilon$

Separate conditional statements: sufficiency alone does not imply an accurate predictor.

Counterexample 1

Equal CLS does not imply equal patch tokens or information loss.

Counterexample 2

Perfect view consistency can choose the same wrong action.

Counterexample 3

Finite queried tails do not bound the best of ALL legal tails.

Bridge selection: a uniform cost error bounds opportunity only over the identical declared tail support.

These conditional statements are Planning Compiler analysis, not control theorems from OSFAL.

Finite noncollision is not generalization; latent L2 is not physical cost; diagnostics are not closed-loop proof.